رواد مصر الرقمية

Ministry of Communications and Information Technology

SALES FORECASTING AND DEMAND PREDICTION

Supervised by: Eng. Mahmoud Talaat

Youth co.

GROUP: GIZ2_AIS4_S9

MAY 8, 2025

Team Members

Abdallah Gasem	Youssef Ali	Mazen Karam
Mostafa Ahmed	Micheal Emad	Abdallah Sayed

Contents

Abstract

This report presents a comprehensive data science project focused on developing a predictive system for sales forecasting and demand prediction. Utilizing historical sales data and external variables such as promotions, holidays, and weather conditions, the project aims to enhance business operations by reducing stockouts and overstocking, optimizing inventory management, and enabling data-driven decision-making across marketing, supply chain, and financial planning. The methodology includes exploratory data analysis, feature engineering, model training with XGBoost, and deployment via a Flask API. Key achievements include an R^2 score of 0.87, identification of seasonal trends, and significant business impacts through improved operational efficiency.

1 Introduction

The ability to accurately forecast sales and predict demand is critical for businesses aiming to optimize their operations and remain competitive. Inefficient inventory management can lead to stockouts, causing lost sales, or overstocking, resulting in increased holding costs and waste. This project, titled ”Sales Forecasting and Demand Prediction,” leverages data science techniques to build a robust predictive system that uses historical sales data and external factors to forecast future demand. The objective is to enable businesses to make informed decisions, reduce operational inefficiencies, and align marketing, supply chain, and financial strategies with real demand.

The report is structured as follows: Section 2 outlines the problem statement, Section 3 describes the proposed solution, Section 4 details the methodology, and subsequent sections cover data collection, preprocessing, exploratory data analysis, model development, deployment, key achievements, business impact, and a conclusion. Each section provides in-depth insights into the processes and outcomes of the project.

2 Problem Statement

2.1 Business Objective

The primary business objective is to improve the accuracy of sales forecasting to optimize inventory levels, reduce waste, and align business operations with actual demand. Accurate forecasting enables businesses to maintain optimal stock levels, ensuring products are available to meet customer demand without incurring unnecessary storage costs.

2.2 Challenges

Traditional forecasting methods often fail to account for complex patterns such as seasonality, time-based trends, and external influences like promotional events, holidays, or weather conditions. These limitations can lead to inaccurate predictions, resulting in overstocking or stockouts. Additionally, customer segment behavior varies significantly, requiring a nuanced approach to forecasting that considers segment-specific trends and preferences.

2.3 Scope

The project focuses on developing a predictive system that incorporates historical sales data and external variables to generate accurate demand forecasts. The system aims to support data-driven decision-making across multiple business functions, including inventory management, marketing strategy, and supply chain planning.

3 Proposed Solution

To address the challenges outlined in the problem statement, we developed a machine learning-based solution that forecasts future sales demand using historical data and time-based features. The proposed solution includes the following key components:

- Segment Analysis through Exploratory Data Analysis (EDA): Analyzing customer segments to identify trends, seasonality, and high-performing groups.
- Feature Engineering: Creating features such as lag variables, promotional indicators, and seasonality markers to capture demand patterns.
- Model Training: Utilizing the XGBoost algorithm to train a high-accuracy predictive model.
- Model Deployment: Deploying the model via a Flask API to provide real-time forecasts for business applications.
- Forecast Generation: Generating actionable predictions to guide inventory management and marketing decisions.

This solution enables businesses to achieve accurate, accessible, and data-driven sales planning, ultimately improving operational efficiency and reducing costs.

4 Methodology

The methodology for the sales forecasting and demand prediction project is structured into six key phases, each contributing to the development and deployment of the predictive system.

4.1 Problem Understanding

The first phase involved defining the business objective and understanding the challenges associated with sales forecasting. The objective was to enhance forecasting accuracy to optimize inventory and planning processes. Key challenges included accounting for time-based patterns and external influences, which traditional methods often overlook.

4.2 Data Collection

Data was sourced from a Kaggle CSV dataset containing historical sales transactions. The dataset included features such as:

- Row ID, Order ID, Order Date, Ship Date

- Ship Mode, Customer ID, Customer Name, Segment
- Country, City, State, Postal Code, Region
- Product ID, Category, Sub-Category, Product Name
- Sales, Quantity, Discount, Profit

The dataset spanned multiple years, capturing transactions across various customer segments and product categories, allowing for comprehensive trend and seasonality analysis.

4.3 Data Preprocessing

Data preprocessing was critical to ensure the quality and usability of the dataset. Key steps included:

- Handling Missing Values: Imputing or removing missing data to maintain dataset integrity.
- Outlier Detection and Treatment: Identifying and addressing outliers to prevent model bias.
- Feature Engineering: Creating time-based features (e.g., Year, Month, Day, Weekday) and lag features (e.g., previous month's sales) to capture temporal patterns.
- Data Aggregation: Grouping sales data by month and segment to create a time-series format suitable for forecasting.

4.4 Exploratory Data Analysis (EDA)

EDA was conducted to uncover insights into sales trends, seasonality, and customer behavior. Key findings included:

- Seasonal Trends: Sales exhibited clear seasonal patterns, with peaks during specific months (e.g., holiday seasons).
- Customer Segment Behavior: Certain segments (e.g., corporate clients) showed higher sales volumes and repeat purchases.
- Shipping Mode Preferences: Most customers preferred cost-effective Standard Class shipping, with faster options like Same Day rarely used.

Visualizations, such as bar charts comparing unique orders to customers and horizontal bar charts showing shipping mode distribution, provided actionable insights into customer behavior.

4.5 Model Development

Multiple machine learning models were tested, including linear regression, decision trees, and gradient boosting algorithms. The XGBoost model was selected due to its superior performance, achieving an R^2 score of 0.87 on test data. Key aspects of model development included:

- Feature Selection: Incorporating engineered features like lags, promotional indicators, and seasonality markers.
- Model Evaluation: Using metrics such as Root Mean Squared Error (RMSE) and R^2 to assess performance.
- Hyperparameter Tuning: Optimizing XGBoost parameters to enhance accuracy and prevent overfitting.

4.6 Model Deployment

The trained XGBoost model was deployed using a Flask API hosted on IBM Cloud, enabling real-time demand predictions. The deployment process involved:

- API Development: Creating endpoints to accept input data and return forecasts.
- Data Validation: Ensuring input data accuracy and compatibility with the model.
- Scalability: Designing the API to handle multiple requests efficiently.

The deployed model provides businesses with accessible, real-time predictions to support inventory and marketing decisions.

5 Data Collection

The data collection phase involved sourcing a comprehensive dataset from Kaggle, which included detailed sales transactions. The dataset's features provided a rich foundation for analysis, covering customer demographics, product details, and transaction specifics. The multi-year span of the data enabled the identification of long-term trends and seasonal patterns, critical for accurate forecasting.

Additional external data, such as promotional events, holidays, and weather conditions, was incorporated to enhance the model's ability to capture external influences on demand. The combination of internal sales data and external variables ensured a holistic approach to forecasting.

6 Data Preprocessing

Data preprocessing was a multi-step process aimed at preparing the dataset for analysis and modeling. The following subsections detail the key activities.

6.1 Handling Missing Values and Outliers

Missing values were addressed through imputation (e.g., using mean or median for numerical features) or removal, depending on the extent of missing data. Outliers were identified using statistical methods (e.g., Z-scores) and treated by capping or removing extreme values to prevent skewing the model.

6.2 Feature Engineering

Feature engineering was a critical component of preprocessing, as it transformed raw data into meaningful inputs for the model. Key features included:

- Time-Based Features: Extracted from Order Date, including Year, Month, Day, and Weekday, to capture temporal patterns.
- Sales Aggregation: Monthly sales grouped by customer segment to create a time-series dataset.
- Lag Features: Previous month's sales and other lag-based features to capture demand momentum.
- External Variables: Indicators for promotions, holidays, and weather conditions to account for external influences.

These features enhanced the model's ability to detect complex patterns and improve forecasting accuracy.

7 Exploratory Data Analysis

EDA provided critical insights into the dataset, guiding feature engineering and model development. Key analyses included:

7.1 Graph 1: Number of Unique Orders vs. Customers

A bar chart comparing unique orders (1,764) to unique customers (707) revealed that customers placed an average of 2.5 orders, indicating repeat business or bulk purchasing behavior. This insight highlighted the importance of modeling customer loyalty and segment-specific purchasing patterns.

7.2 Graph 2: Shipping Mode Distribution

A horizontal bar chart showed that Standard Class shipping was the most popular, while Same Day shipping was rarely used. This finding suggested that customers prioritized cost over speed, informing inventory and logistics planning.

7.3 Seasonality and Segment Analysis

EDA identified strong seasonal trends, with sales peaks during holiday seasons and specific months. Segment-wise analysis revealed that corporate and consumer segments exhibited distinct purchasing behaviors, with corporate clients showing higher sales volumes. These insights informed the inclusion of seasonality and segment-specific features in the model.

8 Model Development

The model development phase involved testing and selecting the best-performing algorithm for sales forecasting. The XGBoost model was chosen due to its ability to handle complex, non-linear relationships and its high accuracy ($R^2 = 0.87$).

8.1 Model Selection

Multiple models were evaluated, including:

- Linear Regression: Simple but unable to capture non-linear patterns.
- Decision Trees: Improved performance but prone to overfitting.
- Random Forest: Better than decision trees but slower and less accurate than XGBoost.
- XGBoost: Best performance due to gradient boosting and robust handling of features.

8.2 Feature Importance

Feature importance analysis revealed that lag features, seasonality markers, and promotional indicators were the most influential in predicting sales. These findings validated the feature engineering approach and guided model optimization.

8.3 Performance Metrics

The model was evaluated using RMSE and R^2. The XGBoost model achieved an R^2 score of 0.87, indicating that 87% of the variance in sales was explained by the model. The low RMSE further confirmed the model's accuracy.

9 Model Deployment

The deployment phase focused on making the model accessible to business users through a Flask API hosted on IBM Cloud. The API was designed to:

- Accept input data in a structured format (e.g., JSON).
- Validate input data to ensure compatibility with the model.
- Return real-time sales forecasts for use in inventory and marketing planning.

The deployment process included rigorous testing to ensure reliability and scalability, allowing the API to handle multiple requests efficiently. The user-friendly interface enabled non-technical users to leverage the model's predictions effectively.

10 Key Achievements

The project achieved several significant milestones, including:

- High Model Accuracy: Developed an XGBoost model with an R^2 score of 0.87, ensuring reliable sales forecasts.
- Insightful EDA: Identified critical seasonal trends, customer behaviors, and high-performing segments.
- Advanced Feature Engineering: Created time-based and lag features to enhance model performance.
- Effective Deployment: Deployed a Flask API for real-time predictions, improving accessibility for business users.

These achievements underscore the project's success in delivering a robust and actionable forecasting solution.

11 Business Impact

The sales forecasting system delivers significant business value by:

- Reducing Stockouts and Overstocking: Accurate demand predictions ensure optimal inventory levels, minimizing lost sales and excess inventory costs.
- Improving Operational Efficiency: Data-driven insights enable better planning for marketing campaigns, staffing, and supply chain logistics.
- Enhancing Decision-Making: Real-time forecasts support strategic decisions across multiple business functions.

By aligning operations with predicted demand, businesses can achieve cost savings, improve customer satisfaction, and maintain a competitive edge.

12 Conclusion

The Sales Forecasting and Demand Prediction project successfully developed a machine learning-based system to forecast future sales demand with high accuracy. By leveraging historical sales data, external variables, and advanced techniques like XGBoost and feature engineering, the project addressed key challenges in inventory management and operational planning. The deployed Flask API ensures that businesses can access real-time predictions, enabling data-driven decision-making. The project's achievements, including an R^2 score of 0.87 and significant business impacts, demonstrate its value in optimizing operations and enhancing competitiveness. Future work could explore additional external variables or advanced deep learning models to further improve accuracy.